

- docs/optimizer.md

KL-bounded Natural-Gradient Updates (closed forms and step size)

State space. Open simplex $\Delta_+^{n-1} = \{p \in \mathbb{R}^n \mid p_i > 0, \sum_i p_i = 1\}$.

KPIs (unit bet).

$$\text{RTP}(p) = \sum_i p_i r_i, \quad \text{Hit}(p) = \sum_{r_i > 0} p_i, \quad \text{Var}(p) = \sum_i p_i r_i^2 - \text{RTP}(p)^2$$

1) One natural-gradient step (information-geometry form)

Gradient of Var w.r.t. p :

$$g_i := \frac{\partial \text{Var}}{\partial p_i} = r_i^2 - 2 \text{RTP}(p) r_i. \quad (2)$$

Center it to lie in the simplex tangent space:

$$\tilde{g}_i := g_i - \langle g \rangle_p, \quad \langle g \rangle_p = \sum_i p_i g_i. \quad (3)$$

A KL/mirror-descent (natural-gradient flavor) update is the **exponential tilt**:

$$q_i(\eta) \propto p_i \exp(\eta \tilde{g}_i), \quad q_i(\eta) = \frac{p_i e^{\eta \tilde{g}_i}}{Z(\eta)}, \quad Z(\eta) = \sum_j p_j e^{\eta \tilde{g}_j}. \quad (4)$$

Properties: $q_i > 0$ (interior point), $\sum_i q_i = 1$ (mass conserved).

2) KL trust-region step size (1-D monotone solve)

We bound the move by $D_{\text{KL}}(q(\eta) \parallel p) \leq \varepsilon$.

For $q(\eta)$ above,

$$D_{\text{KL}}(q(\eta) \parallel p) = \sum_i q_i \log \frac{q_i}{p_i} = \eta \mathbb{E}_{q(\eta)}[\tilde{g}] - \log Z(\eta). \quad (5)$$

Let $A(\eta) := \log Z(\eta)$. Then

$$\mathbb{E}_{q(\eta)}[\tilde{g}] = A'(\eta), \quad D_{\text{KL}}(\eta) = \eta A'(\eta) - A(\eta). \quad (6)$$

Facts: $D_{\text{KL}}(0) = 0$, $D'_{\text{KL}}(\eta) = \eta \text{Var}_{q(\eta)}(\tilde{g}) \geq 0$.

Hence there is a unique $\eta^* > 0$ with $D_{\text{KL}}(\eta^*) = \varepsilon$.

Practical solve (1-D).

$$\text{find } \eta^* \text{ s.t. } f(\eta) := \eta A'(\eta) - A(\eta) - \varepsilon = 0. \quad (7)$$

Use bisection (robust) or Newton:

$$A'(\eta) = \frac{\sum_i p_i \tilde{g}_i e^{\eta \tilde{g}_i}}{Z(\eta)}, \quad A''(\eta) = \text{Var}_{q(\eta)}(\tilde{g}), \quad f'(\eta) = \eta A''(\eta). \quad (8)$$

Newton step: $\eta \leftarrow \eta - \frac{f(\eta)}{f'(\eta)}$ with safeguarding.

3) Iterate and accept

One iteration.

- (i) From $p^{(t)}$ compute RTP and \tilde{g} .
- (ii) Solve $D_{\text{KL}}(\eta) = \varepsilon$ to get $q(\eta)$.
- (iii) If guardrails are violated, project q back (see `constraints.md`).
- (iv) **Accept if**

$$\text{Var}(p^{(t+1)}) > \text{Var}(p^{(t)}) \wedge D_{\text{KL}}(p^{(t+1)} \| p^{(t)}) \leq \varepsilon \wedge \text{GuardrailOK}. \quad (9)$$

Else shrink ε (or η) and retry.

4) Expected improvement (2nd-order audit note)

For small η ,

$$\Delta \text{Var} \approx \eta \langle \tilde{g}, \tilde{g} \rangle_p - \frac{1}{2} \eta^2 \tilde{g}^\top H \tilde{g}, \quad (10)$$

where H is the Hessian of Var restricted to the tangent space.

This supports an **accept rule** and can be logged as an improvement certificate.